

Concurrency

Massimo Merro

4 December 2017

Introduction

Our focus so far has been on the semantics of **sequential computations**. However, many interesting systems are not sequential!

- hardware is intrinsically parallel
- multiple-processor machines
- multi-threading (even on a single processor)
- networked machines
- cyber-physical systems
- IoT devices
- in general, concurrency can increase program performance by scheduling parallel independent tasks on multicore hardware, and can enable responsive user interfaces.

Challenges in concurrent systems

- the state-space of our systems become *larger*, with the *combinatorial explosion*; with n threads, each of which can be in only 2 states, the system has 2^n states!
- the state-space is not only larger but also more complex
- parallel components sharing resources should access them in *mutual exclusion*. If this is not done properly those components may suffer *deadlock* or *starvation*
- computations become *nondeterministic* (unless synchrony is imposed), as different threads operate at different speeds
- concurrency in programming might induce severe problems such as *data races*, i.e., concurrent access to shared data by different threads, with consequent unpredictable or erroneous behavior.

More challenges

- **partial failures** (of some process, of some device in a network, or some persistent storage device); need **transaction mechanisms**
- **communication** between different **environments** with different local resources (e.g. different local stores, or libraries); need consistency mechanisms;
- **communication** between administrative domains with **partial trust** (or, indeed **not trust** at all); protection against malicious attack
- **dealing with contingent complexity** (embedded historical accidents, etc).

On next slides

Theme: as for sequential languages seen up to now, but much more so.

Concurrent languages are a complicated world.

Aim of this lecture: just to give you a taste of a how relatively simple semantics can be used to express some of the fine distinctions. Primarily

- 1 to boost your intuition on reasoning on concurrent systems
- 2 this can support rigorous proofs about crypto systems, cache-coherency protocols, database transactions, etc.



Our Goal: Define the simplest possible concurrent language and explore a few interesting issues.

Parallel composition: Our Design Choices

- threads don't return a value
- threads are anonymous, i.e. they don't have an identity
- termination of a thread cannot be directly observed withing a program
- processes, in general, are given by a pool of concurrent threads
- threads can't be killed externally.



Changes: Typing and operational semantics

$$(T\text{-sq1}) \frac{\Gamma \vdash e_1 : \text{unit} \quad \Gamma \vdash e_2 : \text{unit}}{\Gamma \vdash e_1; e_2 : \text{unit}} \quad (T\text{-sq2}) \frac{\Gamma \vdash e_1 : \text{unit} \quad \Gamma \vdash e_2 : \text{proc}}{\Gamma \vdash e_1; e_2 : \text{proc}}$$

 single thread
 multithread

L'interprete introduce non determinismo perché non so chi può andare avanti, perché dipende da questioni di basso livello

$$(T\text{-par}) \frac{\Gamma \vdash e_1 : T_1 \quad \Gamma \vdash e_2 : T_2}{\Gamma \vdash e_1 \parallel e_2 : \text{proc}} \quad T_1, T_2 \in \{\text{unit}, \text{proc}\}$$

 il programma diventa multithread
 multithread

$$(\text{par-L}) \frac{\langle e_1, s \rangle \rightarrow \langle e'_1, s' \rangle}{\langle e_1 \parallel e_2, s \rangle \rightarrow \langle e'_1 \parallel e_2, s' \rangle} \quad (\text{par-R}) \frac{\langle e_2, s \rangle \rightarrow \langle e'_2, s' \rangle}{\langle e_1 \parallel e_2, s \rangle \rightarrow \langle e_1 \parallel e'_2, s' \rangle}$$

$$(\text{end-L}) \frac{-}{\langle \text{skip} \parallel e, s \rangle \rightarrow \langle e, s \rangle} \quad (\text{end-R}) \frac{-}{\langle e \parallel \text{skip}, s \rangle \rightarrow \langle e, s \rangle}$$

- $\Gamma \vdash e : \text{unit}$ entails e singlethreaded
- $\Gamma \vdash e : \text{proc}$ entails e multithreaded

|| non determinismo permette di gestire sistemi complessi perciò se è gestito non è assolutamente un male

As in any concurrent language:

- threads execute asynchronously - the semantics allows any interleaving of the reductions of the threads
- all threads can read and write the shared memory
- As a consequence, the **Determinacy property does not hold**.

For instance:

$$\langle l := 1 \parallel l := 2, \{l \mapsto 0\} \rangle \rightarrow \langle skip \parallel l := 2, \{l \mapsto 1\} \rangle \rightarrow \langle skip \parallel skip, \{l \mapsto 2\} \rangle \rightarrow \langle skip, \{l \mapsto 2\} \rangle$$

But also

$$\langle l := 1 \parallel l := 2, \{l \mapsto 0\} \rangle \rightarrow \langle l := 1 \parallel skip, \{l \mapsto 2\} \rangle \rightarrow \langle skip \parallel skip, \{l \mapsto 1\} \rangle \rightarrow \langle skip, \{l \mapsto 1\} \rangle$$

Race conditions

- both 'assignments' and 'dereferencing' are **atomic** operations: in the previous configuration we can get a store where location 1 is associated to either 1 or 2. No strange combinations of them.
- However, in $(l := e) \parallel e'$ the semantic steps which are necessary to evaluate e and e' can be interleaved
- So, what about the execution of program $(l := 1 + !l) \parallel (l := 7 + !l)$?
- In this case, we can get **race conditions**, i.e. the output can be something completely unexpected and inconsistent with the intentions of any thread!
- In particular, as depicted at pag. 97 of the notes, there are 3 possible final configurations for $\langle (l := 1 + !l) \parallel (l := 7 + !l), \{l \mapsto 0\} \rangle$:
 - 1 $\langle skip, \{l \mapsto 1\} \rangle$
 - 2 $\langle skip, \{l \mapsto 7\} \rangle$
 - 3 $\langle skip, \{l \mapsto 8\} \rangle$
- (1) and (2) are due to "interferences" while executing the assignments; only (3) corresponds to some correct scheduling!

Morals

- There are too many possible results
- Actually, all the possible executions give rise to a combinatorial explosion of states
- Drawing state-space diagrams, as done at pag. 97 of Sewell's notes, works only for very little examples: we need better techniques to analyze our concurrent programs!
- Almost certainly you (as the programmer) didn't want all those 3 outcomes to be possible - need better idioms or constructs for programming.

How do we get anything coherent done?

- need some way(s) to synchronize between threads, so can enforce **mutual exclusion** for shared data
- Think of Lamport's "Bakery" algorithm for concurrent and distributed systems. Can you code that in our small concurrent language? If not, what would you need in the language?
- though you can depend on built-in support from the scheduler, e.g. **mutexes** or **condition variable** (or, at the lower level, **tas**, test-and-set, or **cas**, compare-and-set).

Adding primitives mutexes in the language

Mutex names $m \in \mathbb{M} = \{m, m_1, \dots\}$

Configurations $\langle e, s, \mu \rangle$, where $\mu : \mathbb{M} \rightarrow \mathbb{B}$ is the mutex state

Expressions $e \in \text{Exp} \dots \mid e \parallel e \mid \text{lock } m \mid \text{unlock } m$

Typing:

(T-lock) $\frac{}{\Gamma \vdash \text{lock } m : \text{unit}}$

(T-unlock) $\frac{}{\Gamma \vdash \text{unlock } m : \text{unit}}$

Operational semantics:

(lock) $\frac{}{\langle \text{lock } m, s, \mu \rangle \rightarrow \langle \text{skip}, s, \mu[m \mapsto \text{true}] \rangle} \text{ if } \neg \mu(m)$

(unlock) $\frac{}{\langle \text{unlock } m, s, \mu \rangle \rightarrow \langle \text{skip}, s, \mu[m \mapsto \text{false}] \rangle}$

... and adapt all the other rules to extended configurations $\langle e, s, \mu \rangle$.

Using a Mutex

To avoid race conditions, we can rewrite the previous program as follows:

$$Prg \stackrel{\text{def}}{=} (\text{lock } m; l := 1 + !l; \text{unlock } m) \parallel (\text{lock } m; l := 7 + !l; \text{unlock } m)$$

Let $\langle Prg, s_0, \mu_0 \rangle$ be a configuration such that $s_0 = \{l \mapsto 0\}$ and μ_0 returns *false* for any mutex name. Then, for all possible executions traces of the configuration $\langle Prg, s_0, \mu_0 \rangle$ we will always have

$$\langle Prg, s_0, \mu_0 \rangle \rightarrow^* \langle \text{skip}, \{l \mapsto 8\}, \mu_0 \rangle$$

No other final configurations are possible!

The two assignments will be executed one after the other in mutual exclusion.

Note that the two assignments *commute*, so we end up in the same final state whichever got the lock first.

Deadlocks

The construct `lock m` can block if the the mutex `m` has already been locked by another thread. So, if we use (at least) two mutexes we can easily **deadlock!**

Consider

$$e = \begin{array}{l} (\text{lock } m_1; \text{lock } m_2; l_1 := !l_2; \text{unlock } m_1; \text{unlock } m_2) \\ \parallel \\ (\text{lock } m_2; \text{lock } m_1; l_2 := !l_1; \text{unlock } m_2; \text{unlock } m_1) \end{array}$$

... **and we don't want deadlocks!**

Language Properties

- Obviously, we don't have **Determinacy** anymore
- **Type preservation** is still valid
- **Typing** and type inferences is scarcely changed
- Very fancy type systems can be used to enforce locking disciplines
- **Progress** in general is not valid unless we adopt a type system to enforce a locking discipline. In that case, we would have deadlock-freedom for free. This has an influence on our notions of semantic equivalence.

Semantic equivalences on concurrent programs

Since deadlocking processes are not ruled out anymore by our type system, we have to revisit our semantic equivalences

Let's amend the typed equivalences seen for sequential computations.

Trace equivalence \simeq_{Γ}

$e_1 \simeq_{\Gamma} e_2$ iff for all mutex states μ and all stores s , s.t. $\text{dom}(\Gamma) \subseteq \text{dom}(s)$, we have $\Gamma \vdash e_1 : T_1$, $\Gamma \vdash e_2 : T_2$, $T_1, T_2 \in \{\text{unit}, \text{proc}\}$, and

- $\langle e_1, s, \mu \rangle \rightarrow^* \langle e'_1, s', \mu' \rangle$ implies $\exists e'_2. \langle e_2, s, \mu \rangle \rightarrow^* \langle e'_2, s', \mu' \rangle$
- $\langle e_2, s, \mu \rangle \rightarrow^* \langle e'_2, s', \mu' \rangle$ implies $\exists e'_1. \langle e_1, s, \mu \rangle \rightarrow^* \langle e'_1, s', \mu' \rangle$.

Notice that now we consider also **partial traces** and not only those leading to final configurations.

Example (1)

$$\begin{aligned}P_1 &= ((\text{lock } m; l := 3) \parallel (\text{lock } m; l := 4)) \\Q_1 &= \text{lock } m; (l := 3 \parallel l := 4)\end{aligned}$$

- Is $P_1 \simeq_{\Gamma} Q_1$, for $\Gamma = \{l : \text{intref}\}$? Yes, it is!
- Is $C[P_1] \simeq_{\Gamma'} C[Q_1]$, for any well-typed context $C[\cdot]$? No, it isn't!
- Consider the context $C[\cdot]$ defined as follows:

$$[\cdot] \parallel (x_1 := !l; x_2 := !l; \text{if } !x_2 = !x_1 + 1 \text{ then } r := 1 \text{ else } r := 0)$$

- Then $\langle C[Q_1], s, \mu \rangle \rightarrow^* \langle \text{skip}, s', \mu' \rangle$, with $s(l)=s(r)=0$, $\mu(m)=\text{false}$, $s'(l) = 4$, $s'(r)=1$ and $\mu'(m)=\text{true}$.
- But, there is no trace $\langle C[P_1], s, \mu \rangle \rightarrow^* \langle \dots, s'', \mu'' \rangle$, with $s' = s''$ and $\mu' = \mu''$! This is because P_1 can “touch” location l only once!
- However, $C[\cdot]$ is not “fair” because it does not acquire the mutex before accessing location l . Any “fair” distinguishing context?

Example (2)

Suppose we can type the following programs:

$$P_2 = ((\text{lock } m; l := 3; \text{unlock } m) \parallel (\text{lock } m; l := 4; \text{unlock } m))$$
$$Q_2 = \text{lock } m; (l := 3 \parallel \text{unlock } m; \text{lock } m \parallel l := 4); \text{unlock } m$$
$$R_2 = \text{lock } m; (l := 3 \parallel \text{unlock } m \parallel \text{lock } m \parallel l := 4); \text{unlock } m$$

In these 3 programs the critical assignments to l are fully locked.

- Is $P_2 \simeq_{\Gamma} Q_2 \simeq_{\Gamma} R_2$, for $\Gamma = \{l_0 : \text{intref}, l : \text{intref}\}$? Yes, it is.
- Is $C[P_2] \simeq_{\Gamma} C[Q_2]$, for any “fair” context $C[\cdot]$? No, it isn't!
- Consider the “fair” context $C[\cdot]$ defined as follows:

$$[\cdot] \parallel (\text{lock } m; x_1 := !l; x_2 := !l; (\text{if } !x_2 = !x_1 + 1 \text{ then } r := 1 \text{ else } r := 0); \text{unlock } m)$$

- Then $\langle C[Q_2], s, \mu \rangle \rightarrow^* \langle \dots, s', \mu' \rangle$, with $s(l) = s(r) = 0$, $\mu(m) = \text{false}$, $s'(l) = 4$, $s'(r) = 1$ and $\mu'(m) = \text{true}$.
- But there is no trace s.t. $\langle C[P_2], s, \mu \rangle \rightarrow^* \langle \dots, s', \mu' \rangle$.

So...

- It is not considering “fair” contexts that we fix the problem!
- What is the problem? \simeq_{τ} is not preserved by parallel contexts!
- Why? Because trace equivalence forgets about intermediate states!

Moral: Parallel contexts have a stronger distinguishing power because they have more chances to create interferences.

- That’s why it is much more difficult to write correct concurrent programs: when you add parallel threads a correct sequential program may go wrong!

Trace Congruence: a much finer semantic equivalence

Trace congruence \cong_Γ

Define $e_1 \cong_\Gamma e_2$ to hold iff for all mutex states μ and all stores s , such that $\text{dom}(\Gamma) \subseteq \text{dom}(s)$, we have $\Gamma \vdash e_1 : T_1$, $\Gamma \vdash e_2 : T_2$, $T_1, T_2 \in \{\text{unit}, \text{proc}\}$ and

- $\langle e_1, s, \mu \rangle \rightarrow^* \langle e'_1, s', \mu' \rangle$ implies $\exists e'_2. \langle e_2, s, \mu \rangle \rightarrow^* \langle e'_2, s', \mu' \rangle$
- $\langle e_2, s, \mu \rangle \rightarrow^* \langle e'_2, s', \mu' \rangle$ implies $\exists e'_1. \langle e_1, s, \mu \rangle \rightarrow^* \langle e'_1, s', \mu' \rangle$
- if $e_1 \cong_\Gamma e_2$ then for any expression e such that $\Gamma' \vdash e_1 \parallel e : \text{proc}$ and $\Gamma' \vdash e_2 \parallel e : \text{proc}$, for some Γ' , then $e_1 \parallel e \cong_{\Gamma'} e_2 \parallel e$.

By definition, the relation \cong_Γ is preserved by parallel contexts!

... and

$$P_1 \not\cong_\Gamma Q_1$$

$$P_2 \not\cong_\Gamma Q_2$$

What about bi-similarity for concurrent programs?

We adapt the definitions to the current setting with mutexes:

Similarity


We say that e_1 is simulated by e_2 , written $e_1 \sqsubseteq_{\Gamma} e_2$, iff

- $\Gamma \vdash e_1 : T_1$ and $\Gamma \vdash e_2 : T_2$, with $T_1, T_2 \in \{\text{unit}, \text{proc}\}$
- for any μ and s with $\text{dom}(\Gamma) \subseteq \text{dom}(s)$, if $\langle e_1, s, \mu \rangle \rightarrow \langle e'_1, s', \mu' \rangle$ then there is e'_2 such that $\langle e_2, s, \mu \rangle \rightarrow^* \langle e'_2, s', \mu' \rangle$, with $e'_1 \sqsubseteq_{\Gamma} e'_2$.


Bisimilarity *— è più forte*

We say that e_1 is bisimilar to e_2 , written $e_1 \approx_{\Gamma} e_2$, iff

- $\Gamma \vdash e_1 : T_1$ and $\Gamma \vdash e_2 : T_2$, with $T_1, T_2 \in \{\text{unit}, \text{proc}\}$
- for any μ and s with $\text{dom}(\Gamma) \subseteq \text{dom}(s)$, if $\langle e_1, s, \mu \rangle \rightarrow \langle e'_1, s', \mu' \rangle$ then there is e'_2 such that $\langle e_2, s, \mu \rangle \rightarrow^* \langle e'_2, s', \mu' \rangle$, with $e'_1 \approx_{\Gamma} e'_2$
- for any μ and s with $\text{dom}(\Gamma) \subseteq \text{dom}(s)$, if $\langle e_2, s, \mu \rangle \rightarrow \langle e'_2, s', \mu' \rangle$ then there is e'_1 such that $\langle e_1, s, \mu \rangle \rightarrow^* \langle e'_1, s', \mu' \rangle$, with $e'_1 \approx_{\Gamma} e'_2$.



$$G[P] \not\approx_t G[Q]$$



$$G[P] \approx G[Q]$$

$$P \approx Q \Rightarrow \begin{cases} P \subseteq Q \\ \wedge \\ P \supseteq Q \end{cases}$$

Now, if you consider the processes of the previous examples:

$$P_1 \sqsubseteq_{\Gamma} Q_1$$

$$Q_1 \not\sqsubseteq_{\Gamma} P_1$$

$$P_2 \sqsubseteq_{\Gamma} Q_2$$

$$Q_2 \not\sqsubseteq_{\Gamma} P_2$$

$$Q_2 \approx_{\Gamma} R_2$$

Unlike \simeq_{Γ} , both \sqsubseteq_{Γ} and \approx_{Γ} can observe changes at **intermediate states!**

In general, $P \sqsubseteq_{\Gamma} Q$ and $Q \sqsubseteq_{\Gamma} P$ does not imply $P \approx_{\Gamma} Q$! Can you find two programs P and Q where this happens?

Is the relation \sqsubseteq_{Γ} a congruence? **Yes, it is!**

Why? Because \sqsubseteq_{Γ} is much sharper when observing processes.

On the power of bi-similarity

Similarity and Bisimilarity are preserved by parallel contexts

- If $e_1 \sqsubseteq_{\Gamma} e_2$ then *for any expression e , such that $\Gamma' \vdash e_1 \parallel e : \text{proc}$ and $\Gamma' \vdash e_2 \parallel e : \text{proc}$, for some Γ' , it holds that*

$$e_1 \parallel e \sqsubseteq_{\Gamma} e_2 \parallel e .$$

- If $e_1 \approx_{\Gamma} e_2$ then *for any expression e , such that $\Gamma' \vdash e_1 \parallel e : \text{proc}$ and $\Gamma' \vdash e_2 \parallel e : \text{proc}$, for some Γ' , it holds that*

$$e_1 \parallel e \approx_{\Gamma} e_2 \parallel e .$$

$$M \parallel A \subseteq M$$

attacco tollerato
non influenza il
sistema

Conditional critical regions

- We have seen that communication between parallel threads is via the store
- In concurrent programs it is very difficult to limit interferences on it
- Many real concurrent programming languages have constructs for alleviating these problems: *semaphores*, *locks*, *critical regions*, etc
- We have seen how to enrich our language with a simple form of locks
- Here we examine a higher-level construct for **conditional critical regions**

↪ *costrutto di alto livello che vuole permettere l'esecuzione di un processo e_2 in maniera atomica*

await e_1 protect e_2 end

- The intuition is that this command may only be executed when the boolean expression e_1 is true, and the entire command e_2 is to be executed to completion without interruption or interference.

- For example consider the program:

$$Prg_1 \stackrel{\text{def}}{=} l := 0 \parallel (\text{await } !l = 0 \text{ protect } l := 1; l := !l + 1 \text{ end})$$

- This is a deterministic program; if it is executed in a state s , with $s(l) \neq 0$, then it will terminate and the only possible terminal state is $s[l \mapsto 2]$.
- As another example consider the more involved program:

$$Prg_2 \stackrel{\text{def}}{=} \begin{array}{l} (\text{await true protect } l_1 := 1; l_1 := !l_0 + 1 \text{ end}) \\ \parallel \\ (\text{await true protect } l_0 := 2; l_0 := !l_1 + 1 \text{ end}) \end{array}$$

- The two guards, set to true, are vacuous, so which protected command is executed first is chosen non-deterministically.

Formally...

await false do skip; e₁

Language:

Configurations $\langle e, s \rangle$, as before

Expressions $e \in \text{Exp} ::= \dots \mid e \parallel e \mid \text{await } e \text{ protect } e \text{ end}$

Typing:

$$\text{(T-await)} \quad \frac{\Gamma \vdash e_1 : \text{bool} \quad \Gamma \vdash e_2 : \text{unit}}{\Gamma \vdash \text{await } e_1 \text{ protect } e_2 \text{ end} : \text{unit}}$$

Operational semantics:

$$\text{(await)} \quad \frac{\langle e_1, s \rangle \rightarrow^* \langle \text{true}, s' \rangle \quad \langle e_2, s' \rangle \rightarrow^* \langle \text{skip}, s'' \rangle}{\langle \text{await } e_1 \text{ protect } e_2 \text{ end}, s \rangle \rightarrow \langle \text{skip}, s'' \rangle}$$

This is a kind of **test-and-set command**: whenever the guard e_1 evaluates to true the command e_2 can be executed **atomically**, in just one step!

It is easy to see that

await *true* protect ($l := !l + 1; l := !l - 1$) end

\approx_{Γ}

await *false* protect ($l := !l + 1; l := !l - 1$) end

\approx_{Γ}

await *e* protect ($l := !l + 1; l := !l - 1$) end

\approx_{Γ}

skip

for any expression *e* such that

- $\Gamma \vdash e : \text{bool}$
- *e* does not modify the store.

Example

Let us consider the following programs:

$$\begin{aligned} P_4 &\stackrel{\text{def}}{=} l_0 := 0; \\ &\quad (\text{await } !l_0 = 0 \text{ protect } (l := 1; l_0 := 1) \text{ end}) \\ &\quad || \\ &\quad (\text{await } !l_0 = 0 \text{ protect } (l := 0; l_0 := 1) \text{ end}) \end{aligned}$$
$$Q_4 \stackrel{\text{def}}{=} l_0 := 0; (l := 0; l_0 := 1 \parallel l := 1; l_0 := 1)$$

Supponendo di poter tipare il seguente processo:

$$R_4 \stackrel{\text{def}}{=} l_0 := 0; (\text{await } !l_0 = 0 \text{ protect } (l := 0; l_0 := 1 \parallel l := 1; l_0 := 1) \text{ end})$$

- $P_4 \sqsubseteq_{\Gamma} Q_4$ *P_4 simula Q_4*
- $Q_4 \not\sqsubseteq_{\Gamma} P_4$ *Q_4 non simula P_4*
- $P_4 \approx_{\Gamma} R_4$ *bisimili*

Nondeterministic choice

Let us suppose to enrich our language with the following construct:

Configurations $\langle e, s \rangle$, as before

Expressions $e \in \text{Exp} ::= \dots \mid e + e$

scelta non
deterministica

Typing:

$$\text{(T-choice)} \quad \frac{\Gamma \vdash e_1 : \text{unit} \quad \Gamma \vdash e_2 : \text{unit}}{\Gamma \vdash e_1 + e_2 : \text{unit}}$$

Operational semantics:

$$\text{(ChoiceL)} \quad \frac{\langle e_1, s \rangle \rightarrow \langle e'_1, s' \rangle}{\langle e_1 + e_2, s \rangle \rightarrow \langle e'_1, s' \rangle}$$

e_1 prende il
sopraelevato, e_2
viene scartato


$$\text{(ChoiceR)} \quad \frac{\langle e_2, s \rangle \rightarrow \langle e'_2, s' \rangle}{\langle e_1 + e_2, s \rangle \rightarrow \langle e'_2, s' \rangle}$$

e_2 prende il sopraelevato
 e_1 viene scartato

This construct chooses nondeterministically one branch or the other; once a branch is chosen the other one is discarded!

skip + $l:=1, l:=2 \xrightarrow{l:=1} l:=2$

qui non vado mai in skip
perché è un processo morto mentre
l'altro può tranquillamente proseguire



in questo caso e con questa
implementazione non avrei il medesimo
comportamento

↳ poiché devo valutare se entrambi i
processi possono proseguire, in quel
caso tiro la moneta

skip + skip $\xrightarrow{\text{deadlock}}$

skip $\equiv_R P$

skip è simulato da
qualsiasi processo
basta che dia
fermo tutto che non
fa nulla

True and false algebraic laws

- ① $e + e \overset{\text{bisimile}}{\approx}_{\Gamma} e$
- ② $e \approx_{\Gamma} \text{skip}; e$
- ③ $e_1 + e_2 \sqsubseteq_{\Gamma} (\text{skip}; e_1) + e_2$
- ④ $e_1 + e_2 \sqsupseteq_{\Gamma} (\text{skip}; e_1) + e_2$
- ⑤ $e_1 + e_2 \not\approx_{\Gamma} (\text{skip}; e_1) + e_2$
- ⑥ $e_1 + e_2 \sqsubseteq_{\Gamma} (\text{skip}; e_1) + (\text{skip}; e_2)$
- ⑦ $e_1 + e_2 \sqsupseteq_{\Gamma} (\text{skip}; e_1) + (\text{skip}; e_2)$
- ⑧ $e_1 + e_2 \not\approx_{\Gamma} (\text{skip}; e_1) + (\text{skip}; e_2)$
- ⑨ $e + e \approx_{\Gamma} (\text{skip}; e) + e$
- ⑩ $e + e \approx_{\Gamma} (\text{skip}; e) + (\text{skip}; e)$

$$① \quad e + e \stackrel{?}{\approx} e$$

$$\frac{e_1 \rightarrow e'_1}{e_1 + e_2 \rightarrow e'_1}$$

$$\frac{e_2 \rightarrow e'_2}{e_1 + e_2 \rightarrow e'_2}$$

$$e + e \stackrel{?}{\approx} e$$

$$\downarrow$$

$$e'$$

$$\downarrow$$

$$e'$$

$$e \rightarrow e'$$

perché $e \rightarrow e'$

$$② \quad e \stackrel{?}{\approx} \text{skip}; e$$

$$\downarrow$$

$$e'$$

$$\downarrow$$

$$e$$

$$\downarrow$$

$$e'$$

bisimulazione dice che si può rispondere con uno o più passi

Allora $e \sqsubseteq_r \text{skip}; e$

$$e \stackrel{?}{\approx} \text{skip}; e$$

$$\downarrow$$

$$e$$

$$\downarrow$$

$$e$$

$$\downarrow$$

$$e$$

$$e \sqsubseteq \text{skip}; e$$

$$\rightarrow$$

$$H_0$$

$$e \approx_r \text{skip}; e$$

$$\textcircled{3} \quad e_1 + e_2 \equiv \overbrace{\text{skip}; e_1 + e_2}^{P \quad Q}$$

Provo a costruire una relazione di simulazione $R = \{(P, Q), \dots\}$

① Analizziamo i passi possibili nella 1^a coppia (da sx vs dx)

1.1 La $P = e_1 + e_2 \rightarrow e_1' \equiv P'$ $e_1 \rightarrow e_1' \equiv P'$ $e_2 \rightarrow e_2'$ e ho applicato la (choice L). Come risponde Q?

Allora $Q \equiv \text{skip}; e_1 + e_2 \rightarrow e_1 \rightarrow e_1' \equiv Q'$ con $(P', Q') \in R$

1.2 La $P = e_1 + e_2 \rightarrow e_2' \equiv P'$ perché $e_2 \rightarrow e_2'$ e ho applicato la (choice R). Cosa risponde Q?

Allora $Q \equiv \text{skip}; e_1 + e_2 \rightarrow e_2' \equiv Q'$, con $(P', Q') \in R$

$$\textcircled{4} \quad e_1 + e_2 \equiv \overbrace{\text{skip}; e_1 + e_2}^{P \quad Q} \quad R = \{(P, Q), \dots\} \cup Id$$

① Analizziamo la 1^a coppia

1) La $Q \equiv \text{skip}; e_1 + e_2 \rightarrow e_1 \equiv P'$, per una applicazione delle regole. P può simulare Q ?


Allora $P \rightarrow P' \equiv P'$ con $(Q', P') \in R$

2) La $Q \equiv \text{skip}; e_1 + e_2 \rightarrow e_2'$, applicando la (choice R). Ma allora $Q = e_1 + e_2 \rightarrow e_2' \equiv Q$, applicando la (choice R) con $(Q', P') \in Id \subseteq R$

$$\begin{aligned} e_1 &\equiv l := 1 \\ e_2 &\equiv l := 2 \end{aligned}$$

② Analizziamo la 2^a coppia

2.1 La $e_1 \rightarrow e_1'$, per qualche e_2' . Allora $e_1 + e_2 \rightarrow e_1'$, applicando la (choice L) con $(e_1', e_2') \in Id \subseteq R$



$$R = \{(q, p). (e_1; e_2)\} \cup Id$$

$$e_1 \equiv l := 5$$

$$e_2 \equiv m := 3$$

$$e_1 + e_2 \stackrel{?}{=} skip;e_1 + skip;e_2$$

Q

1) Analizziamo 1^a coppia di R

1.1) Sia $Q \rightarrow e_1$, per b (choice L). Allora $P \rightarrow^* P = e_1 + e_2$, con $(e_1, e_1 + e_2) \in R$

1.2) Sia $Q \rightarrow e_2$, per b (choice R). Allora $P \rightarrow^* P = e_1 + e_2$, con $(e_2, e_1 + e_2) \in R$

2) Analizziamo la 2^a coppia

La $e_1 \rightarrow e_1$, ma allora $e_1 + e_2 \rightarrow e_1'$ (choice L) con $(e_1', e_1) \in Id \subseteq R$

3) Analizziamo la terza coppia

finisce nel caso precedente

$$e_1 + e_2 \stackrel{?}{=} skip;e_1 + skip;e_2$$

Q

non genera altre coppie
in R perché non
ricade nell'identif

$$R = \{(p, q). \quad \quad \quad \} \cup Id$$

1) Analizziamo 1^a coppia di R

Example: When execution order is important

Consider the following algebraic law:

$$l := 1 \parallel m := 2 \approx_{\Gamma} (l := 1; m := 2) + (m := 2; l := 1)$$

Can we generalise this law as follows?

$$e_1 \parallel e_2 \approx_{\Gamma} (e_1; e_2) + (e_2; e_1)$$

for arbitrary expressions e_1 and e_2 ?

$$e_1 \quad l = 0; m = 2$$

$$e_2 \quad l = 3; m = 4$$

*Non vale in generale ma solo su
certe condizioni*

$$e_1 \equiv l=0, l:=l+1$$

$$e_2 \equiv l=1$$

$$e_1 \parallel e_2 \stackrel{?}{\approx} e_1; e_2 + e_2; e_1$$

$$\downarrow l=0$$

$$l:=l+1 \parallel l:=1$$

$$\downarrow l=1$$

$$l:=l+1$$

$$\downarrow$$

$$l=2$$

$$\downarrow l=0$$

$$l:=l+1; l:=1$$

$$\downarrow l=1$$

$$l=1$$

Si rompe la simulazione (E)
l'altro resto valido

Example: On Bisimulation

Let

- $e_1 \stackrel{\text{def}}{=} (l := 1) + (l := 1; m := 2)$

- $e_2 \stackrel{\text{def}}{=} l := 1; m := 2$


which of the following statements is true?

- $e_1 \sqsubseteq_{\Gamma} e_2$
- $e_1 \sqsupseteq_{\Gamma} e_2$

se valgono le simulazioni
in entrambi i versi non
sempre vale la bisimulazione

- $e_1 \approx_{\Gamma} e_2$.

↓
se vale la
bisimulazione
sicuramente
valgono le simulazioni
in entrambi i versi



An encoding of nondeterministic choice

Question: Is $e_1 + e_2$ a primitive construct or it can be codified?

An encoding of nondeterministic choice

Question: Is $e_1 + e_2$ a primitive construct or it can be codified?

Let us try to encode nondeterministic choice using parallel composition, locations and the construct for critical regions:

$$e_1 \uplus e_2 \stackrel{\text{def}}{=} \text{let } m : \text{ref int} = \text{ref } 0 \text{ in}$$

\downarrow
macro

$\begin{aligned} &(\text{await } !m = 0 \text{ protect } m := 1 \text{ end}; e_1 \\ &\parallel \\ &\text{await } !m = 0 \text{ protect } m := 1 \text{ end}; e_2) \end{aligned}$

Said in other words: does our **implementation** of nondeterministic choice satisfy its **specification**, or... something close to it? Actually:

- $e_1 \uplus e_2 \sqsupseteq_{\Gamma} e_1 + e_2$
- $e_1 \uplus e_2 \sqsubseteq_{\Gamma} e_1 + e_2$
- $e_1 \uplus e_2 \not\approx_{\Gamma} e_1 + e_2$
- $e_1 \uplus e_2 \approx_{\Gamma} (\text{skip}; e_1) + (\text{skip}; e_2).$

Persistent behaviours (1)

We know that when we write $e_1; e_2$ we have to execute e_1 first, and only when e_1 has been completed we can execute e_2 .

However, how can we write in our language a program that repeats subsequently the same program e ?

$e; e; e; e; \dots$

Proposal:

$$\begin{aligned} \text{RepSeq}(e) &\stackrel{\text{def}}{=} \text{let } S : (\text{unit} \rightarrow \text{unit}) \rightarrow (\text{unit} \rightarrow \text{unit}) \\ &= (\text{fn } f : \text{unit} \rightarrow \text{unit} \Rightarrow (\text{fn } x : \text{unit} \Rightarrow x; (f\ x))) \\ &\quad \text{in fix.} S\ e \end{aligned}$$

Suppose to have a CBN semantics!

Persistent behaviours (2)

What about a program that forks an arbitrary number of threads e ?

$$e \parallel e \parallel e \parallel e \parallel \dots$$

Proposal:

$$\begin{aligned} \text{RepPar}(e) &\stackrel{\text{def}}{=} \text{let } P : (\text{unit} \rightarrow \text{unit}) \rightarrow (\text{unit} \rightarrow \text{unit}) \\ &= (\text{fn } f : \text{unit} \rightarrow \text{unit} \Rightarrow (\text{fn } x : \text{unit} \Rightarrow x \parallel (f\ x))) \\ &\quad \text{in fix.} P\ e \end{aligned}$$

Again suppose to have a CBN semantics!

Data race and critical regions (1)

During the execution of the program $\text{RepSeq}(l := !l + 1)$ the value associated to the location l increases monotonically:

$$l := !l + 1; l := !l + 1; l := !l + 1; \dots$$

Whereas during the execution of the program $\text{RepPar}(l := !l + 1)$ the value associated to the location l may increase or decrease.

$$l := !l + 1 \parallel l := !l + 1 \parallel l := !l + 1 \parallel \dots$$

This is because this program suffers of **data races** at locations l .

Actually:

- $\text{RepSeq}(l := !l + 1) \sqsubseteq_{\Gamma} \text{RepPar}(l := !l + 1)$
- $\text{RepSeq}(l := !l + 1) \not\sqsubseteq_{\Gamma} \text{RepPar}(l := !l + 1)$

↗ può anche scegliere di eseguire uno per volta

Data races and critical regions (2)

Any way to avoid those data races maintaining concurrency?

Proposal:

$$\begin{aligned} \text{AwtPar}(e) &\stackrel{\text{def}}{=} \\ &\text{let } A : (\text{unit} \rightarrow \text{unit}) \rightarrow (\text{unit} \rightarrow \text{unit}) \\ &= (\text{fn } f : \text{unit} \rightarrow \text{unit} \Rightarrow (\text{fn } x : \text{unit} \Rightarrow \text{await } \text{true} \text{ protect } x \text{ end } \parallel (f\ x))) \\ &\text{in fix}.A\ e \end{aligned}$$

Now, it is possible to prove that

$$\text{RepSeq}(l := !l + 1) \approx_{\Gamma} \text{AwtPar}(l := !l + 1).$$