

# Laboratorio di Software Security 2023/2024

## Lezione 1

Marco Campion

[marco.campion@univr.it](mailto:marco.campion@univr.it)



UNIVERSITÀ  
di **VERONA**

*Inria*



PSL



# Goals

- 1) Understanding some well known software **vulnerabilities**
- 2) Understanding how to **exploit** such vulnerabilities as an attacker
- 3) Understanding the potential impact of such attacks and how to **prevent** them

# This lab is a starting point!

- This is an **introductory** lab on software attacks/hacking/defenses!
- **Good starting point** for who is interested to study more complex software attacks/defenses
- How to learn **more** on this topic:
  - Cyberchallenge univr (<https://cyberchallenge.it/>)
  - Online tutorials ([codearcana](#), [roman](#), [nightmare](#), ...)
  - CTF contests (<https://ctftime.org/>)
  - ...

# Software vulnerabilities

- A *vulnerability* is a weakness which can be exploited by an *attacker* to perform *unauthorized actions* within your program
- To exploit a vulnerability, an attacker relies on *tools and techniques* related to a *software weakness*

# What is a vulnerability?

- **Question:** Is the following code *vulnerable*?

```
int authenticate() {
    char* password = "MyPassword!";
    char name[10];
    char* psw = malloc(256);

    printf("Enter your name: ");
    scanf("%s", name);
    printf("Enter the password: ");
    scanf("%s", psw);

    if (strcmp(password, psw) == 0) {
        printf("Authenticated with name:\n");
        printf(name);
        return 1;
    } else {
        printf("The password is wrong! Please try again!\n");
        return 0;
    }
}
```

# What is a vulnerability?

- **Question:** Is the following code *vulnerable*? **Yes!**

There are (at least)  
**4 vulnerabilities** in  
this code

```
int authenticate() {
    char* password = "MyPassword!";
    char name[10];
    char* psw = malloc(256);

    printf("Enter your name: ");
    scanf("%s", name);
    printf("Enter the password: ");
    scanf("%s", psw);

    if (strcmp(password, psw) == 0) {
        printf("Authenticated with name:\n");
        printf(name);
        return 1;
    } else {
        printf("The password is wrong! Please try again!\n");
        return 0;
    }
}
```

# What is a vulnerability?

- **Question:** Is the following code *vulnerable*? **Yes!**

There are (at least) 4 vulnerabilities in this code:

- *Hardcoded password*

```
int authenticate() {  
    char* password = "MyPassword!";  
    char name[10];  
    char* psw = malloc(256);  
  
    printf("Enter your name: ");  
    scanf("%s", name);  
    printf("Enter the password: ");  
    scanf("%s", psw);  
  
    if (strcmp(password, psw) == 0) {  
        printf("Authenticated with name:\n");  
        printf(name);  
        return 1;  
    } else {  
        printf("The password is wrong! Please try again!\n");  
        return 0;  
    }  
}
```

# What is a vulnerability?

- **Question:** Is the following code *vulnerable*? **Yes!**

There are (at least) 4 vulnerabilities in this code:

- *Hardcoded password*
- *Potential stack-overflow*

```
int authenticate() {
    char* password = "MyPassword!";
    char name[10];
    char* psw = malloc(256);

    printf("Enter your name: ");
    scanf("%s", name);
    printf("Enter the password: ");
    scanf("%s", psw);

    if (strcmp(password, psw) == 0) {
        printf("Authenticated with name:\n");
        printf(name);
        return 1;
    } else {
        printf("The password is wrong! Please try again!\n");
        return 0;
    }
}
```



# What is a vulnerability?

- **Question:** Is the following code *vulnerable*? **Yes!**

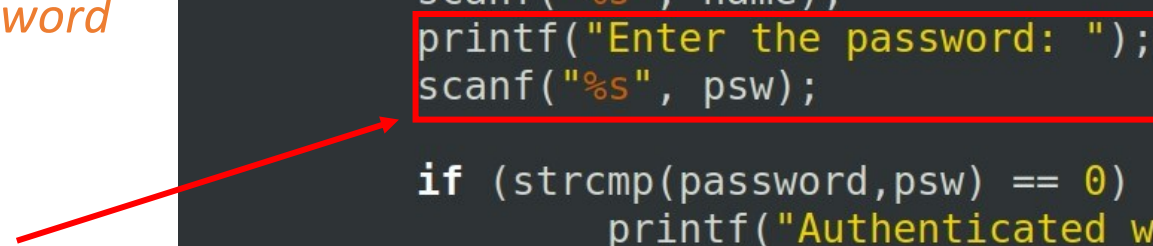
There are (at least) 4 vulnerabilities in this code:

- *Hardcoded password*
- *Potential stack-overflow*
- *Potential heap-overflow*

```
int authenticate() {
    char* password = "MyPassword!";
    char name[10];
    char* psw = malloc(256);

    printf("Enter your name: ");
    scanf("%s", name);
    printf("Enter the password: ");
    scanf("%s", psw);

    if (strcmp(password, psw) == 0) {
        printf("Authenticated with name:\n");
        printf(name);
        return 1;
    } else {
        printf("The password is wrong! Please try again!\n");
        return 0;
    }
}
```



# What is a vulnerability?

- **Question:** Is the following code *vulnerable*? **Yes!**

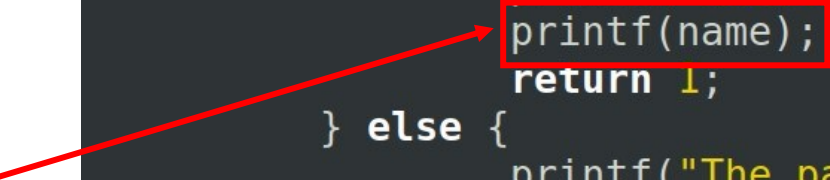
There are (at least) 4 vulnerabilities in this code:

- *Hardcoded password*
- *Potential stack-overflow*
- *Potential heap-overflow*
- *Potential format-string vulnerability*

```
int authenticate() {
    char* password = "MyPassword!";
    char name[10];
    char* psw = malloc(256);

    printf("Enter your name: ");
    scanf("%s", name);
    printf("Enter the password: ");
    scanf("%s", psw);

    if (strcmp(password, psw) == 0) {
        printf("Authenticated with name:\n");
        printf(name);
        return 1;
    } else {
        printf("The password is wrong! Please try again!\n");
        return 0;
    }
}
```



# Sources of vulnerabilities

- Complexity, inadequacy, and (uncontrolled) changes
- Incorrect or changing assumptions (capabilities, inputs, outputs)
- Flawed specifications and designs
- Poor implementation of software interfaces (input validation, error and exception handling)
- Unintended, unexpected interactions with other components with the software's execution environment
- *Inadequate knowledge of secure coding practices*

# Lab Software Security program overview

- Background on x86 architectures
- Debugging with gdb
- Reverse engineering
- Patching executables
- Python 3 library : pwntools
- Buffer overflow attacks: variables overriding and RA corruption
- Buffer overflow attacks: arbitrary code execution
- OS and Compiler-level defenses

# Prerequisites for this Lab

- Basic knowledge of C/C++ and Unix-based OS

# Prerequisites for this Lab

- Basic knowledge of C/C++ and Unix-based OS
- Desktop/laptop with *Intel x86/x86-64* microprocessor and Unix-based OS (e.g. Ubuntu)
  - If **you have** a machine with *Intel x86/x86-64* microprocessor :
    - Either you manually install on your Unix OS all the necessary tools
    - Or you can download our Ubuntu virtual machine with everything you need
  - If **you don't have** a machine with *Intel x86/x86-64* microprocessor :
    - Use Virtual Lab (*new from this year!*)

**More infos at the end of this lecture!**

# Prerequisites for this Lab

- Basic knowledge of C/C++ and Unix-based OS
- Desktop/laptop computer with *Intel x86/x86-64* microprocessor and Unix-based OS (e.g. Ubuntu)
- **Very good knowledge of x86 architectures!**

# In this lecture

- Background on x86 architectures
- Debugging with gdb
- Reverse engineering
- Patching executables
- Python 3 library : pwntools
- Buffer overflow attacks: variables overriding and RA corruption
- Buffer overflow attacks: arbitrary code execution
- OS and Compiler-level defenses



# Roadmap

C:

```
car *c = malloc(sizeof(car));  
c->miles = 100;  
c->gals = 17;  
float mpg = get_mpg(c);  
free(c);
```

Java:

```
Car c = new Car();  
c.setMiles(100);  
c.setGals(17);  
float mpg =  
    c.getMPG();
```

Assembly  
language:

```
get_mpg:  
    pushq   %rbp  
    movq    %rsp, %rbp  
    ...  
    popq    %rbp  
    ret
```


Machine  
code:

```
0111010000011000  
100011010000010000000010  
1000100111000010  
110000011111101000011111
```

OS:



Computer  
system:




# From code to programs

➤ Compiling a C program is a multi-stage process composed of four steps:

- 1) preprocessing
- 2) compilation
- 3) assembly
- 4) linking

# From code to programs: preprocessing

- In the first phase some *preprocessor* commands (in C they start with '#') are interpreted:



# From code to programs: compilation

- In the second phase, preprocessed code is translated into *assembly instructions*:

```
#include <stdio.h>

#define MESSAGE "Hello world!"

int main() {
    printf(MESSAGE);
    return 0;
}
```

gcc -s hello.c

```
# 2 "hello.c" 2

# 5 "hello.c"
int main() {
    printf("Hello world!");
    return 0;
}
```

```
main:
.LFB0:
    .cfi_startproc
    pushq   %rbp
    .cfi_def_cfa_offset 16
    .cfi_offset 6, -16
    movq    %rsp, %rbp
    .cfi_def_cfa_register 6
    movl    $.LC0, %edi
    movl    $0, %eax
    call    printf
    movl    $0, %eax
    popq    %rbp
    .cfi_def_cfa 7, 8
    ret
```

# From code to programs: assembly

- In the *assembly* phase assembly instructions are translated into *machine* or *object code*:

```
#include <stdio.h>

#define MESSAGE "Hello world!"

int main() {
    printf(MESSAGE);
    return 0;
}
```

gcc -c hello.c



hello.o

```
# 2 "hello.c" 2

# 5 "hello.c"
int main() {
    printf("Hello world!");
    return 0;
}
```

```
main:
.LFB0:
    .cfi_startproc
    pushq   %rbp
    .cfi_def_cfa_offset 16
    .cfi_offset 6, -16
    movq    %rsp, %rbp
    .cfi_def_cfa_register 6
    movl    $.LC0, %edi
    movl    $0, %eax
    call    printf
    movl    $0, %eax
    popq    %rbp
    .cfi_def_cfa 7, 8
    ret
```

# From code to programs: linking

- In the last phase (multiple) *object code* are combined in a single executable
- In the generated file, references (links) to the used library are added

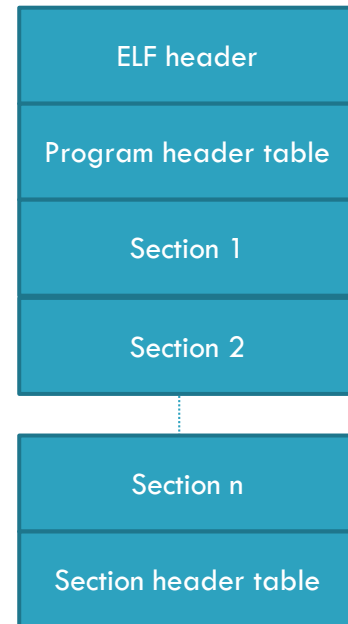


# Static vs Dynamic linking

- Two approaches can be used in the linking phase:
  - **Static Link**
    - Binaries are *self-contained* and do not depend on any external libraries
  - **Dynamic Link**
    - Binaries rely on system libraries that are loaded when needed
    - Mechanisms are needed to *dynamically* relocate code

# Executable and Linkable Format (ELF)

- Standard file format for executables in Unix-like systems
- Any ELF file is structured as:
  - an **ELF header** describing the file content for execution
  - a **Program header table** providing info about how to create a process image
  - a sequence of **Sections** containing what is needed for linking (instructions, data, symbol table, relocation information,...)
  - a **Section header table** with a description of previous sections
- Analogous format in Windows is *Portable Executable* (PE)
- Analogous format in macOS and iOS is *Mach-O*





# ELF: Overview

## DISSECTED FILE

```
~$ uname -m  
armv7l  
~$ ./simple.ARM  
Hello World!
```

### HEADER<sup>1/2</sup>

TECHNICAL DETAILS FOR IDENTIFICATION AND EXECUTION

### SECTIONS

CONTENTS OF THE EXECUTABLE

### HEADER<sup>2/2</sup>

TECHNICAL DETAILS FOR LINKING (IGNORED FOR EXECUTION)

#### ELF HEADER

IDENTIFY AS AN ELF TYPE  
SPECIFY THE ARCHITECTURE

#### PROGRAM HEADER TABLE

EXECUTION INFORMATION

#### CODE

EXECUTABLE INFORMATION

#### DATA

INFORMATION USED BY THE CODE

#### SECTIONS' NAMES

#### SECTION HEADER TABLE

LINKING (CONNECTING PROGRAM OBJECTS) INFORMATION

#### HEXADECMAL DUMP

#### ASCII DUMP

#### FIELDS

#### VALUES

#### EXPLANATION

1	e_ident EI_MAG EI_CLASS, EI_DATA EI_VERSION e_type e_machine e_version e_entry e_phoff e_shoff e_ehsize e_phsize e_phnum e_shsize e_shnum e_shstrndx	0x7F, "ELF" 1 (LITTLE), 1 (LITTLE) 2 (LITTLE) 28 (LITTLE) 1 (LITTLE) 0x80000000 0x40 0x80 0x34 0x20 1 0x20 4 3*	CONSTANT SIGNATURE 32 BITS, LITTLE-ENDIAN ALWAYS 1 EXECUTABLE ARM PROCESSOR ALWAYS 1 ADDRESS WHERE EXECUTION STARTS PROGRAM HEADERS OFFSET SECTION HEADERS OFFSET ELF HEADERS SIZE SIZE OF A SINGLE PROGRAM HEADER COUNT OF PROGRAM HEADERS SIZE OF A SINGLE SECTION HEADER COUNT OF SECTION HEADERS INDEX OF THE NAMES SECTION IN THE TABLE
2	p_type p_offset p_vaddr p_paddr p_filesz p_memsz p_flags	1 (LITTLE) 0 0x80000000 0x80000000 0x90 0x90 5 (LITTLE)	THE SEGMENT SHOULD BE LOADED IN MEMORY OFFSET WHERE IT SHOULD BE READ VIRTUAL ADDRESS WHERE IT SHOULD BE LOADED PHYSICAL ADDRESS WHERE IT SHOULD BE LOADED SIZE ON FILE SIZE IN MEMORY READABLE AND EXECUTABLE

#### ARM ASSEMBLY

#### EQUIVALENT C CODE

3	MOV r2, #13 ADD r1, PC, #20 MOV r8, #1 MOV r7, #1 SVC 0  MOV r8, #1 MOV r7, #1 SVC 0	/* 13 */ /* 20 */ /* 1 */ /* 1 */ /* SVC 0 */  /* 1 */ /* 1 */ /* SVC 0 */	write(STDOUT_FILENO, "hello world!\n", 10);  exit(0);
---	--	--	---

#### STRINGS

"Hello World!\n", 0

#### SECTION NAMES

.shstrtab .text .rodata

#### SECTION HEADER TABLE


INDEX	NAME	TYPE	FLAGS	ADDRESS	OFFSET	SIZE
0	<null>	0	0			
1	.text	1 (LITTLE-ENDIAN)	6 (LITTLE-ENDIAN)	0x80000060	0x60	0x20
2	.rodata	1 (LITTLE-ENDIAN)	2 (LITTLE-ENDIAN)	0x80000080	0x80	0x00
3*	.shstrtab	3 (LITTLE-ENDIAN)	0		0x90	0x19

# ELF: Relevant sections

- **.text**: contains the executable instructions of a program
- **.bss**: contains uninitialised data that contribute to the program's memory image
- **.data, .data1**: contain initialized data that contribute to the program's memory image
- **.rodata, .rodata1**: are similar to **.data** and **.data1**, but refer to read-only data
- **.symtab**: contains the program's symbol table
- **.dynamic**: provides linking information

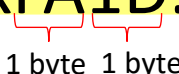
# Recall storage sizes

- A bit can be either 0 or 1
- 1 BYTE = 8 bit, 1 WORD = 2 bytes, 1 DWORD = 4 bytes, 1 QUADWORD = 8 bytes
- The basic storage unit for all data in x86 architectures is 1 byte
- Least significant bit (LSB) = the right-most bit
- Most significant bit (MSB) = the left-most bit



# Hexadecimal integers


- In x86 assembly, hexadecimal (hex) values are used as a compact form for representing binary numbers
  - Base 16 number representation
  - Use characters '0' to '9' and 'A' to 'F'
- Each digit in a hex integer represents 4 bits
- Two hex digits together represent a byte
- In C language they are written as 0xFA1D...  

  
1 byte 1 byte

Decimal	Binary	Hexadecimal
0	0000	0
1	0001	1
2	0010	2
3	0011	3
4	0100	4
5	0101	5
6	0110	6
7	0111	7
8	1000	8
9	1001	9
10	1010	A
11	1011	B
12	1100	C
13	1101	D
14	1110	E
15	1111	F

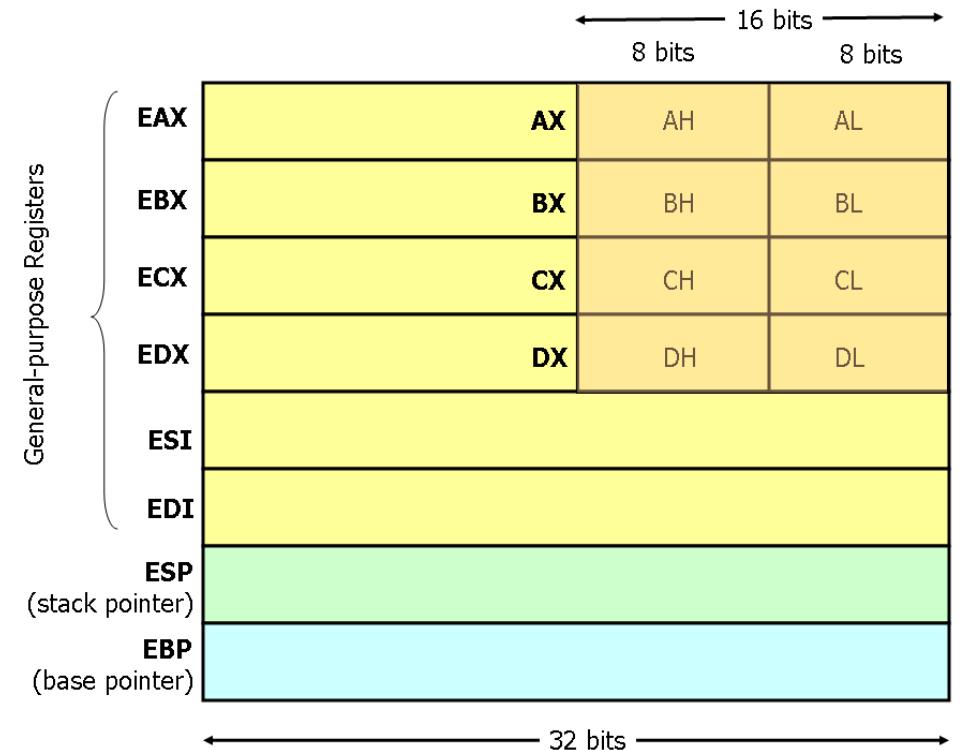
# x86-32 Registers

- x86-32 processors have eight 32-bit general purpose registers
- The register names are mostly historical...
  - EAX used to be called *the accumulator* since it was used by a number of arithmetic operations
  - ECX was known as the *counter* since it was used to hold a loop index
- Two are reserved for special purposes:
  - the *stack pointer* (ESP)
  - the *base pointer* (EBP)



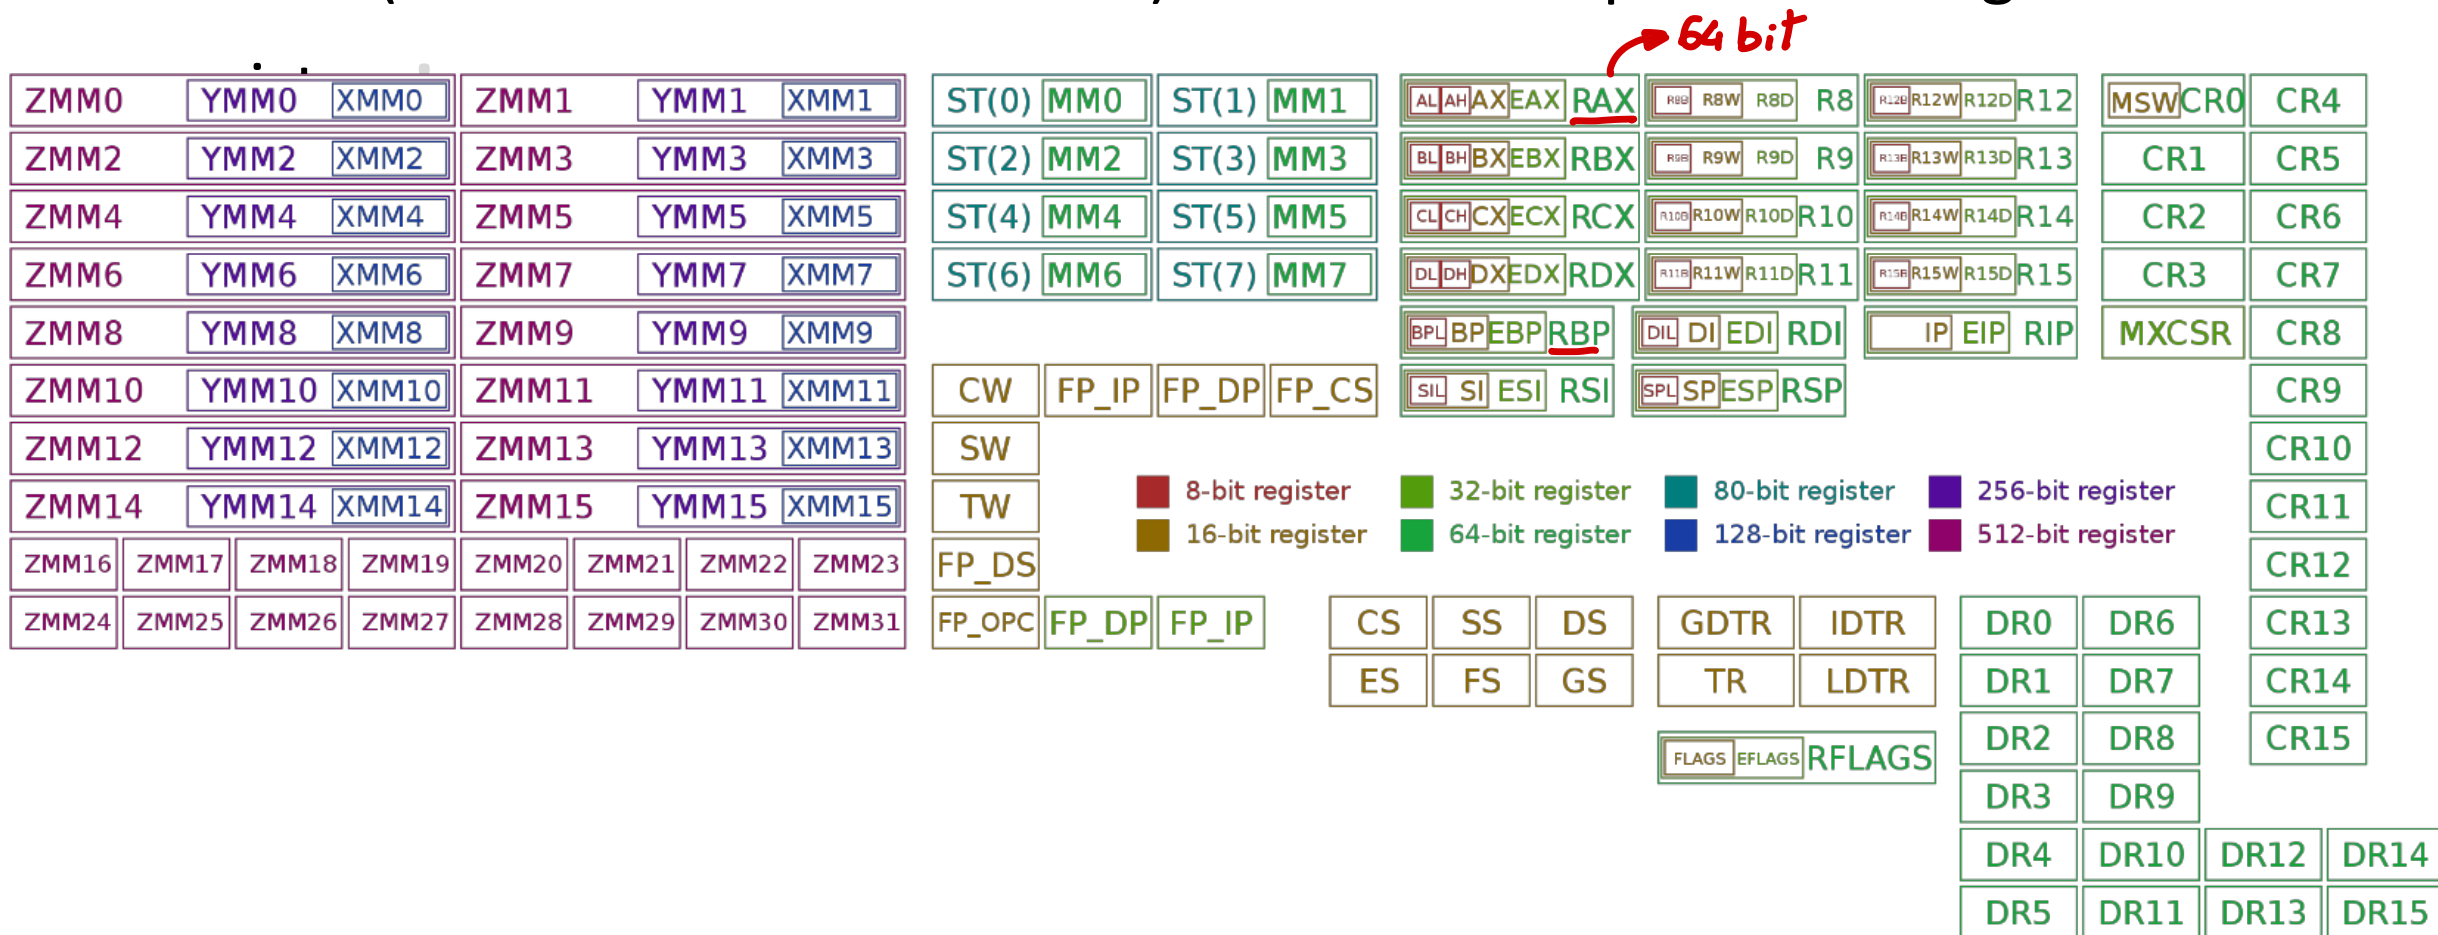
# x86-32 Registers

- For each register, subsections may be used
- For example:
  - AX refers to the least significant 2 bytes of EAX
  - AL refers to the least significant byte of AX
  - AH refers to the most significant byte of AX
- These sub-registers are mainly hold-overs from older, 16-bit versions of the instruction set



# x86-64 Registers

- x86-64 (also called amd64 or x64) architectures provide a larger set of



# x86 Memory management

- In some programming languages, like C, memory management can be controlled by programmers:
  - memory can be **dynamically** allocated and deallocated
  - memory address of variables can be obtained (pointers)
- If  $x$  is a variable,  $\&x$  denotes **the pointer to  $x$** , i.e., the memory address where  $x$  is stored



# Memory allocation

- Let us consider the following simple C program:

```
#include <stdio.h>

int main() {
    int i;
    char c;
    short s;
    long l;

    printf("i is allocated at %p\n", &i);
    printf("c is allocated at %p\n", &c);
    printf("s is allocated at %p\n", &s);
    printf("l is allocated at %p\n", &l);
}
```

Variable declarations


Variable addresses

We can assume that:


- *int* needs 4 bytes
- *char* needs 1 byte
- *short* needs 2 bytes
- *long* needs 8 bytes

# Memory allocation

- Memory is just a sequence (array) of bytes each with a unique address:




- Compilers may introduce **padding** or change the order of data in memory for optimization
- Memory is represented as groups of bytes (matrix) where each row has  $n$  bytes with  $n$  = processor word (e.g. for 64bit architectures,  $n = 8$  bytes)



# Memory allocation: example on a 64bit architecture

## ❖ A “64-bit (8-byte) word-aligned” view of memory:


- In this type of picture, each row is composed of 8 bytes
- Each cell is a byte
- A 64-bit pointer will fit on one row



# Adresses and Pointers

- ❖ An *address* is a location in memory
- ❖ A *pointer* is a data object that holds an address
  - Address can point to *any* data
- ❖ Value 351 stored at address 0x08
  - $351_{10} = 15F_{16}$   
= 0x 00 00 01 5F
- ❖ Pointer stored at 0x38 points to address 0x08

								Address
								0x00
	00	00	00	00	00	00	01 5F	0x08
								0x10
								0x18
								0x20
								0x28
								0x30
	00	00	00	00	00	00	00 08	0x38
								0x40
								0x48



# Sizes of data types in bytes

Java Data Type	C Data Type	32-bit (old)	x86-64
<u>boolean</u>	bool	1	1
<u>byte</u>	char	1	1
<u>char</u>		2	2
<u>short</u>	short int	2	2
<u>int</u>	int	4	4
<u>float</u>	float	4	4
	long int	4	8
<u>double</u>	double	8	8
<u>long</u>	long	8	8
	long double	8	16
<u>(reference)</u>	<b>pointer *</b>	<b>4</b>	<b>8</b>

# Byte ordering

Endianness → convenzione su come  
salvare i byte in memoria


Bigendian  
Little endian

- ❖ How should bytes within a word be ordered in memory?
  - **Example:** store the 4-byte (32-bit) `int`:  
0x a1 b2 c3 d4
- ❖ By convention, ordering of bytes called *endianness*
  - The two options are big-endian and little-endian

# Byte ordering

- ❖ Big-endian (SPARC, z/Architecture)
  - Least significant byte has highest address



- ❖ **Example:** 4-byte data 0xa1b2c3d4 at address 0x100



# Byte ordering


- ❖ **Big-endian** (SPARC, z/Architecture)
  - Least significant byte has highest address
- ❖ **Little-endian** (x86, x86-64)
  - Least significant byte has lowest address

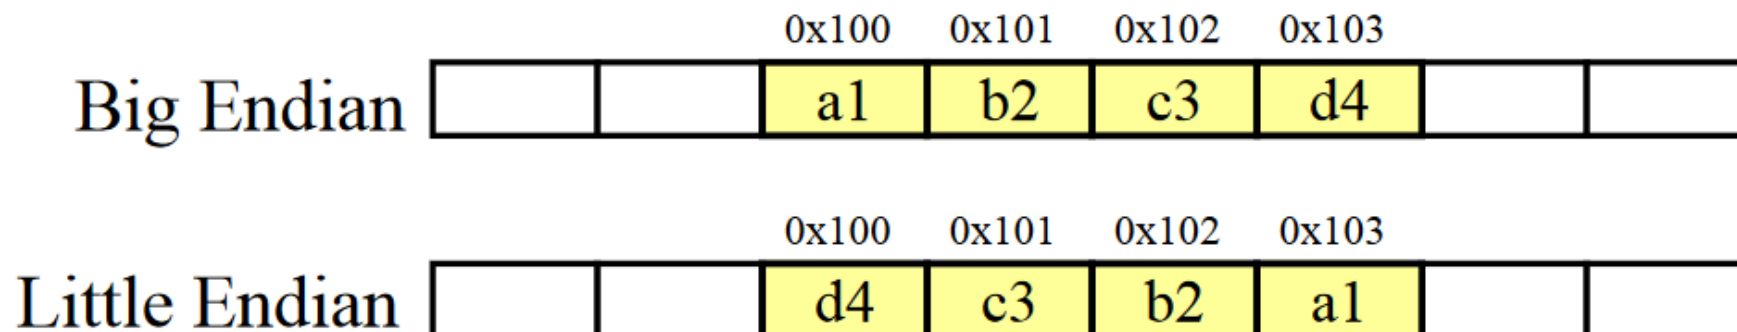
- ❖ **Example:** 4-byte data 0xa1b2c3d4 at address 0x100






# Byte ordering

- ❖ **Big-endian** (SPARC, z/Architecture)
  - Least significant byte has highest address
- ❖ **Little-endian** (x86, x86-64) 
  - Least significant byte has lowest address
- ❖ **Bi-endian** (ARM, PowerPC)
  - Endianness can be specified as big or little
- ❖ **Example:** 4-byte data 0xa1b2c3d4 at address 0x100



# Memory segments

- Memory is allocated for each **process** (a running program) to store **data** and **code**.
- This allocated memory consists of different **segments**:
  - **stack**: for local variables
  - **heap**: for dynamic memory
  - **data segment**:
    - *global uninitialized variables (.bss)*
    - *global initialized variables (.data)*
  - **code segment**



# Intel x86 Instruction Sets

- We provide a short introduction of a small but useful subset of the available instructions and assembler directives of Intel x86 assembly language
- A detailed description can be found at the following links:
  - Short Assembly Guide
    - <https://www.cs.virginia.edu/~evans/cs216/guides/x86.html>
  - Online tutorial
    - [https://www.tutorialspoint.com/assembly\\_programming/index.htm](https://www.tutorialspoint.com/assembly_programming/index.htm)
  - Free e-book
    - <http://mirror.easynome.at/nongnu/pgubook/ProgrammingGroundUp-1-0-booksize.pdf>
  - Intel's Pentium Manuals
    - <http://www.intel.com/content/www/us/en/processors/architectures-software-developer-manuals.html>

# x86 Instructions : Intel vs AT&T syntaxes

- x86 architectures (both 32- and 64-bit) has two alternative syntaxes available for assembly language

➤ AT&T syntax: `movl $1, %eax`

➤ Intel syntax: `mov eax, 1`



**We will consider  
this syntax!**

# x86 Instructions : Data Movement Instructions

## ➤ MOV : copies data from right to left

### *Syntax*

```
mov <reg>, <reg>  
mov <reg>, <mem>  
mov <mem>, <reg>  
mov <reg>, <const>  
mov <mem>, <const>
```

- Square brackets [ val ] are used to directly access memory address contained in val

### *Examples*

```
mov eax, ebx — copy the value in ebx into eax  
mov BYTE PTR [ebx], 2 ; Move 2 into the single byte at the address stored in EBX.
```

# x86 Instructions : Data Movement Instructions

- **PUSH** : places its operand onto the top of the stack

## *Syntax*

`push <reg32>`

`push <mem>`

`push <con32>`

## *Examples*

`push eax` — push `eax` on the stack

`push [var]` — push the 4 bytes at address `var` onto the stack

# x86 Instructions : Data Movement Instructions

- **POP : removes the 4-byte data from the top of the stack**

## *Syntax*

`pop <reg32>`

`pop <mem>`

## *Examples*

`pop edi` — pop the top element of the stack into EDI.

`pop [ebx]` — pop the top element of the stack into memory at the four bytes starting at location EBX.

# x86 Instructions : Data Movement Instructions

## ➤ LEA : load effective address

### *Syntax*

`lea <reg32>, <mem>`

### *Examples*

`lea edi, [ebx+4*esi]` — the quantity `EBX+4*ESI` is placed in EDI.

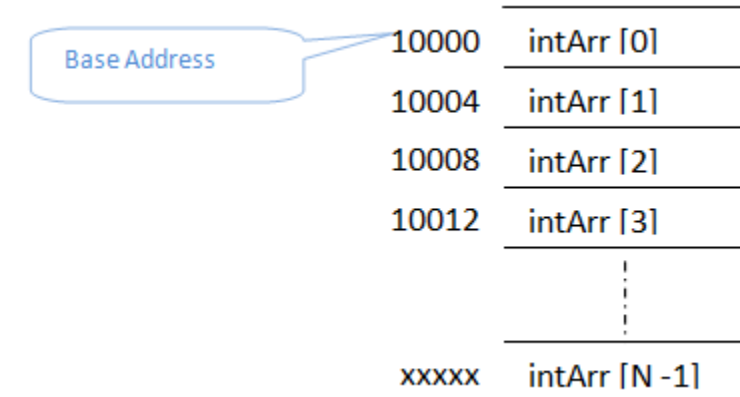
`lea eax, [var]` — the value in *var* is placed in EAX.

`lea eax, [val]` — the value *val* is placed in EAX.



# C arrays in memory

- Each element of a C Array is allocated contiguously
- Initial address of the array = address of the first element
- Each element will occupy the memory space required to accomodate the values for its type, i.e., depending on elements datatype (1, 2, 4 or 8 bytes)
- Total memory allocated to an array = number of elements x size of one element
- MOV and LEA instructions used with arrays



A diagram showing the memory layout of an array. A blue callout box labeled "Base Address" points to the first row of a table. The table has two columns: the first column contains memory addresses (10000, 10004, 10008, 10012, and a gap followed by xxxxx), and the second column contains array element names (intArr [0], intArr [1], intArr [2], intArr [3], and intArr [N -1]). A vertical dashed line is between the last two rows.

10000	intArr [0]
10004	intArr [1]
10008	intArr [2]
10012	intArr [3]
	⋮
xxxxx	intArr [N -1]

`mov eax, [ebx+4*esi]`  
Load in eax the **value** at adress `ebx+4*esi`  
`eax = array[esi]`

vs

`lea eax, [ebx+4*esi]`  
Load in eax the **address** `ebx+4*esi`  
`eax = &array[esi]`

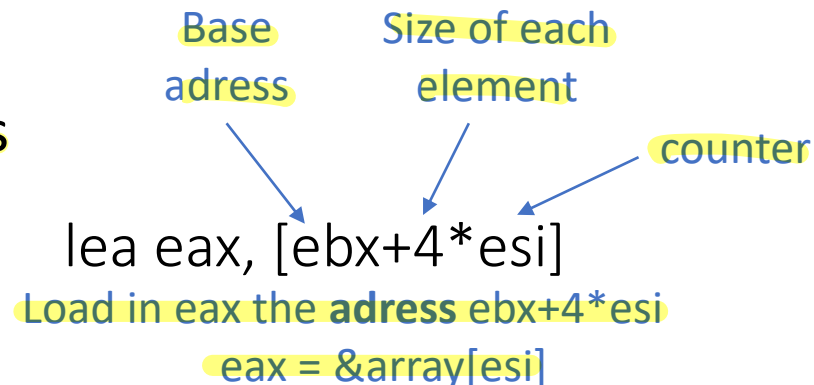


Diagram illustrating the components of the LEA instruction: `lea eax, [ebx+4*esi]`. Three blue arrows point to parts of the instruction: "Base adress" points to `ebx`, "Size of each element" points to `4`, and "counter" points to `esi`.

# C arrays in memory

- The order of the elements of a C array are unaffected by endianness!
  - The first element of an array (array[0]) is always at the lowest address

char s[5] = {'c', 'i', 'a', 'o'}

<b>63</b>	<b>69</b>	<b>61</b>	<b>6F</b>	<b>00</b>
1000	1001	1002	1003	1004

Little-endian  
and big-endian

- <https://www.ascii-code.com/>

- Each element of the array is affected by endianness!

int s[2] = {0x12345678, 0x9ABCDE}

<b>78 56 34 12</b>	<b>DE BC 9A 00</b>
1000	1004

Little-endian

<b>12 34 56 78</b>	<b>00 9A BC DE</b>
1000	1004

Big-endian

# x86 Instructions : Arithmetic and Logic Instructions

## ➤ Arithmetic and Logic Instructions

➤ *add op1,op2*: stores in *op1* the result of *op2+op1*

➤ *sub op1,op2*: stores in *op1* the result of *op2-op1*

➤ *inc op*: increments *op* by one

➤ *and op1,op2*

➤ *or op1,op2*

➤ *xor op1,op2*

➤ ...

} Perform the specified logical operation on the operands, storing the result in the first operand location

# Control Flow

- x86 processor maintains an *Instruction Pointer* (IP) register of 32(or 64)-bit indicating the location in memory where the current instruction starts
- Normally, it increments to point to the next instruction in memory begins after execution an instruction
- Control flow instructions, like jumps, can update IP to point to specific labels

# x86 Instructions: Control Flow Instructions

## ➤ JMP: unconditional jump

`jmp` — Jump

Transfers program control flow to the instruction at the memory location indicated by the operand.

*Syntax*

`jmp <label>`

*Example*

`jmp begin` — Jump to the instruction labeled `begin`.

```
        mov esi, [ebp+8]
begin:  xor ecx, ecx
        mov eax, [esi]
```

# x86 Instructions: Control Flow Instructions

## ➤ Jcondition : conditional jump

### *Syntax*

`je <label>` (jump when equal)  
`jne <label>` (jump when not equal)  
`jz <label>` (jump when last result was zero)  
`jg <label>` (jump when greater than)  
`jge <label>` (jump when greater than or equal to)  
`j1 <label>` (jump when less than)  
`jle <label>` (jump when less than or equal to)

### *Example*

```
cmp eax, ebx
jle done
```

If the contents of EAX are less than or equal to the contents of EBX, jump to the label *done*. Otherwise, continue to the next instruction.

# x86 Instructions: Control Flow Instructions

- **CMP** : compare the values of the two specified operands, setting the condition codes in the machine status word appropriately

## *Syntax*

```
cmp <reg>, <reg>  
cmp <reg>, <mem>  
cmp <mem>, <reg>  
cmp <reg>, <con>
```

## *Example*

```
cmp DWORD PTR [var], 10  
jeq loop
```

If the 4 bytes stored at location *var* are equal to the 4-byte integer constant 10, jump to the location labeled *loop*.

# x86 Instructions: Control Flow Instructions

- **CALL** : pushes the current code location onto the stack and then performs an unconditional jump to the code location indicated by the label operand. It is used to call functions.
- **RET** : implements a return from a function. It first pops a code location off the stack and then performs an unconditional jump to the retrieved code location

## *Syntax*

`call <label>`

`ret`

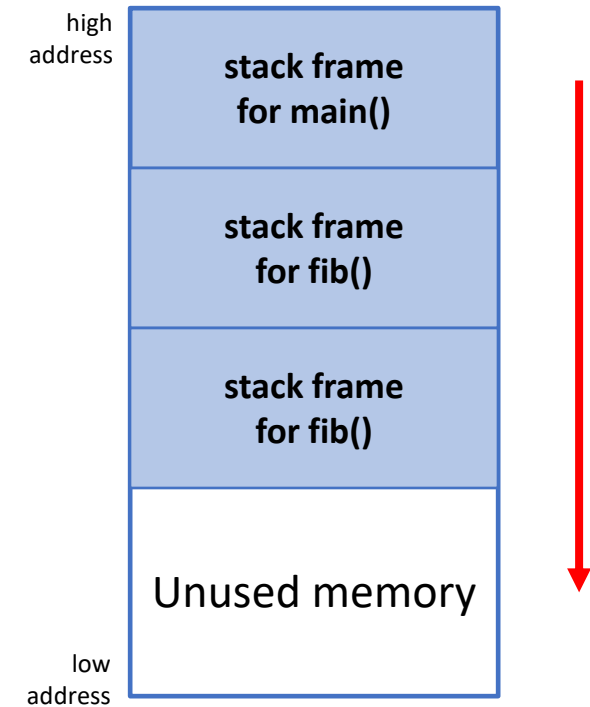


# The stack

diviso in frame logici, cresce verso il basso.


- The stack consists of a sequence of *stack frames* (or activation records), each for each function call:
  - allocated on *call*
  - de-allocated on *return*

```
int main(int argc, char **argv) {  
    int f = fib(n: 10);  
    printf("FIB(10)=%d\n", f);  
}  
  
int fib(int n) {  
    int f1;  
    int f2;  
    if (n<=2) {  
        return 1;  
    } else {  
        f1 = fib(n: n-1);  
        f2 = fib(n: n-2);  
        return f1+f2;  
    }  
}
```



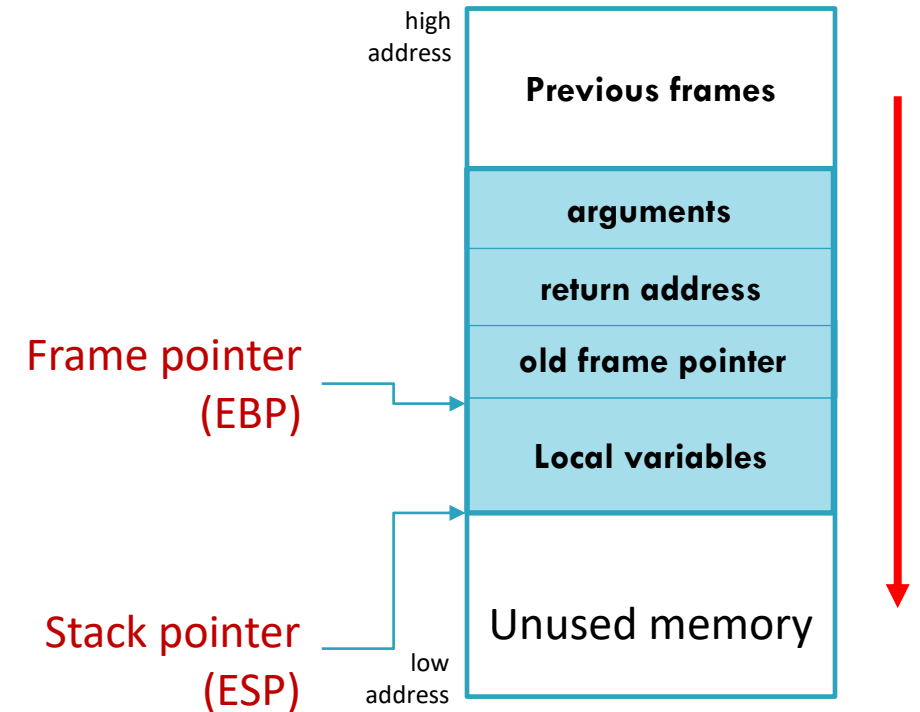
# The stack

- The **stack pointer (SP)** refers to the last element on the stack
- On x86 architectures, the stack pointer is stored in the **ESP (Extended Stack Pointer)** register



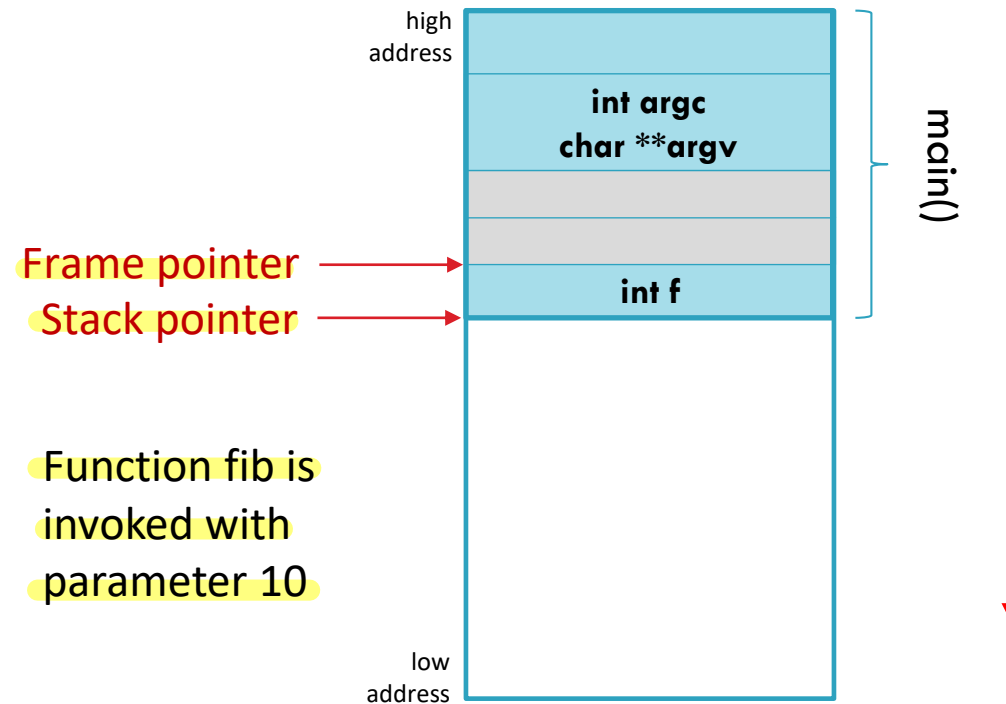
# Stack frame (for x86)

- In x86 architecture, each stack frame contains (in order):
  - Function arguments
  - Copies of registries that must be restored:
    - return address
    - previous frame pointer
  - Local variables
- Frame pointer, named Extended Base Pointer (EBP), provides a starting point to local variables



# Stack example

```
int main(int argc, char **argv) {  
    int f = fib(n: 10);  
    printf("FIB(10)=%d\n", f);  
}  
  
int fib(int n) {  
    int f1;  
    int f2;  
    if (n<=2) {  
        return 1;  
    } else {  
        f1 = fib(n: n-1);  
        f2 = fib(n: n-2);  
        return f1+f2;  
    }  
}
```



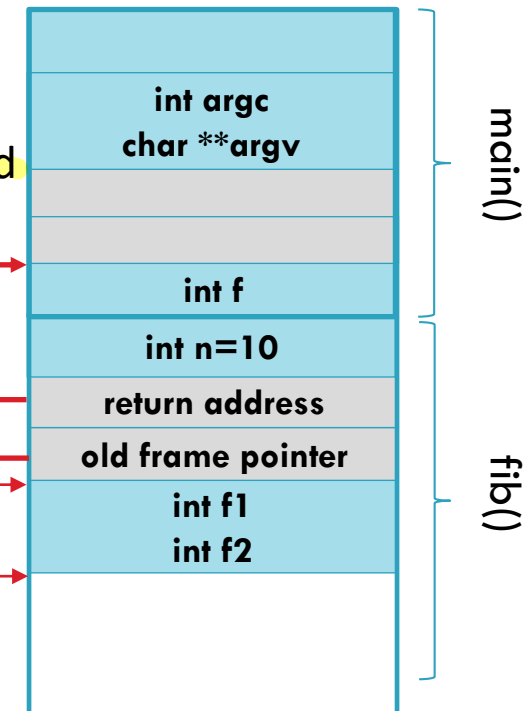
# Stack example

```
int main(int argc, char **argv) {  
    int f = fib(n: 10);  
    printf("FIB(10)=%d\n", f);  
}  
  
int fib(int n) {  
    int f1;  
    int f2;  
    if (n <= 2) {  
        return 1;  
    } else {  
        f1 = fib(n: n-1);  
        f2 = fib(n: n-2);  
        return f1+f2;  
    }  
}
```

Stack frame is  
allocated and  
pointers updated

Frame pointer

Stack pointer




# Stack example

```
int main(int argc, char **argv) {  
    int f = fib(n: 10);  
    printf("FIB(10)=%d\n", f);  
}  
  
int fib(int n) {  
    int f1;  
    int f2;  
    if (n<=2) {  
        return 1;  
    } else {  
        f1 = fib(n: n-1);  
        f2 = fib(n: n-2);  
        return f1+f2;  
    }  
}
```

Frame pointer


Stack pointer

When a function returns, pointers are updated. Function result (if any) is copied in a register



# The heap

- Memory allocation and de-allocation in the stack is very fast
  - However, this memory cannot be used after a function returns
- The heap is used to store dynamically allocated data that outlive function calls:
  - This area is under programmer's responsibility



# Memory management functions

- Basic C functions for memory management are:
  - *malloc(int)*, given an integer *n* allocates an area of *n* (continuous) bytes and returns a **pointer** to that area
  - *free(void\*)*, deallocates the memory associated with a pointer



# Debugging

- A **debugger** is a software tool that allows to:
  - run the target program under controlled conditions
  - track program operations in progress
  - monitor changes in computer resources
  - display the contents of memory
  - modify memory or register contents

# Debugging

- Compilers can be instrumented to emit extra data that debugger can use for a more informative execution
- In the case of gcc compiler, the parameter `-g` can be used
- Debugging information is stored in specific sections of ELF file:
  - `.debug`: contains info for symbolic debugging
  - `.line`: contains line number informations

```
CC> gcc -o hello -g hello.c
CC> readelf --debug-dump hello
Contents of the .debug_aranges section:
```

Length:	44
Version:	2
Offset into .debug_info:	0x0
Pointer Size:	8
Segment Size:	0

Address	Length
0000000000400526	000000000000001a
0000000000000000	0000000000000000

```
Contents of the .debug_info section:
```

```
Compilation Unit @ offset 0x0:
Length:      0x8d (32-bit)
Version:     4
Abbrev Offset: 0x0
```

# GDB : The GNU Project Debugger

- GDB is a debugging tool that can be used to...
  - Start your program, specifying anything that might affect its behavior
  - Stop your program at the occurrence of specified conditions
  - Examine what happened, when your program stopped
  - Change memory or registers content while your program is running

# GDB : The GNU Project Debugger

- GDB supports multiple languages:
  - Assembly (x86, ARM, MIPS...)
  - C
  - C++
  - Rust
  - ...

# GDB : The GNU Project Debugger

- GDB is a terminal tool and can be launched by executing program *gdb*
- You can invoke *gdb* by passing the program to debug

*gdb <program>*

- you can pass the *process ID* as a second argument to debug a running process:

*gdb <program> <pid>*

- Or equivalently use:

*gdb -p <pid>*

# GDB : Startup

When executed, gdb performs the following main steps:

- sets up the command interpreter as specified by the command line
- Load configuration files
- Load symbols of debugged program
- Waits for user commands

# GDB : running and arguments

- A *gdb* command consists of a single line of input, containing:
  - a *command name*
  - a sequence of *parameters*
- Command *run* can be used to *start* a program in *gdb*
- Command *help* can be used to access the list of available commands
- Commands *set args* and *show args* can be used to set and show program arguments

# GDB : Stopping and Counting

- The principal reason to use a debugger is that we can stop a program before it terminates to check its status and, if we experience some problems, investigate and find out why.
- Inside *gdb*, a program may stop for:
  - a *breakpoint*
  - a *signal*
  - the completion of the execution of a *step*



# GDB : Breakpoints

- A *breakpoint* makes your program stop whenever a certain point in the program is reached
  - details can be added to the breakpoint to control in finer detail whether your program stops
- A *watchpoint* is a special breakpoint that stops your program when the value of an expression changes
- A *catchpoint* is another special breakpoint that stops your program when a certain kind of event occurs

# GDB : Breakpoints

- Breakpoints are set with the *break* command (abbreviated *b*):
  - *break location* set break point at the given location (addresses must be preceded by \*)
  - *break* set break point at the next instruction
  - *break [location] if <cond>* set break point with the given condition
- A breakpoint location can be:
  - A line number
  - A label/function name
  - An address (must be preceded by \*)

# GDB : Continuing and Stepping

- When a program stops at a breakpoint, its execution can be resumed exploiting 2 functionalities:
  - Continuing with `continue` or `c` means resuming program execution until another breakpoint is found or your program completes normally
  - Next instruction `nexti` or `ni` means executing just one instruction

# GDB : Inspecting the stack

- When your program has stopped, the first thing we need to know is where it *stopped* and *how* it got there
- The first thing to consider is the content of the *stack*
- gdb commands are available to examine the stack and to read the content of the *stack* and to read any of the stored *stack frames*
- In the *stack* frame you can find:
  - the location of the call in your program
  - the arguments of the call
  - the local variables of the function being called

# GDB :

- The following commands can be used to read the content of the stack:
  - `frame [<selection>]`
    - prints a brief description of the selected stack frame.
  - `info frame [<selection>]`
    - prints a verbose description of the selected stack frame
- See *gdb* documentation for a (long) list of options

# GDB : Other commands

- Command *disas* can be used to disassemble a given function
- The *print* command (abbreviated *p*) prints the content of an address or register (registers must be preceded by \$)  
Example: `p *0x0809aaf6`
- Command *x* treats the content of a memory location or register as an address and prints its content  
Example: `x $eax`
- Command *call* calls a function of our inspected program  
Example: `call (void) function()`
- Command *set* modifies the value of a memory location or a register  
Example: `set $pc=0x400344`

# GEF : GDB Enhanced Features

- GEF consists of a set of commands that extends GDB with additional features for *dynamic analysis* and *exploit development*.
- GEF is based on GDB Python API
- Main GEF features:
  - Embedded hexdump view
  - Automatic dereferencing of *data* and *registers*
  - Heap analysis
  - Display ELF information
- Detailed GEF documentation is available at <https://gef.readthedocs.io/en/master/>

Now...

...hands on!



# Tools for the lab

If **you already have** a computer with Intel processor then, in order to be able to successfully do the exercises, you can either:

a) manually download and install each tool we need:

- gcc-multilib (*sudo apt install gcc-multilib*)
- GEF (<https://gef.readthedocs.io/en/master/>)
- Ghidra (<https://ghidra-sre.org/>)
- Python 3 (<https://www.python.org/downloads/>)
- Pwntools (<https://docs.pwntools.com/en/latest/index.html>)

# Tools for the lab

b) download our Ubuntu virtualbox with everything already installed:

[https://univr-my.sharepoint.com/:u:/g/personal/marco\\_campion\\_univr\\_it/ESoL6ITf6YBMibrti16PfjoBXGQHVmUrKrtM6DxrQGK5IQ?e=jDbyVw](https://univr-my.sharepoint.com/:u:/g/personal/marco_campion_univr_it/ESoL6ITf6YBMibrti16PfjoBXGQHVmUrKrtM6DxrQGK5IQ?e=jDbyVw)

Username : swsec

Password : password

Import it with Oracle Virtual Box (<https://www.virtualbox.org/>)

# Tools for the lab

If **you don't** have a machine with Intel processor then:

- 1) Connect to the univr intranet (either by connecting to the univr wifi or through vpn)
- 2) Visit <https://virtualab.univr.it>
- 3) Use your GIA credentials to authenticate
- 4) Select «*aula\_virtuale*»
- 5) Select any virtual pc and start using the Ubuntu machine!

- Python 3 is already installed
- You can manually install GEF, Ghidra and pwntools with no root permissions
- gcc-multilib may not be installed!

Remember to logout at the end of your work!

# Tutorial Exercise

- Disable ASLR on your system (it will be automatically reactivated on next restart):

```
echo 0 | sudo tee /proc/sys/kernel/randomize_va_space
```

- Run and analyze the program *tutorial* and answer the following questions:

- 1) How many bytes are saved in the stack for the local variables of function *fun()*?
- 2) What are the addresses (in the form register  $\pm$  offset) of the two integer variables of function *fun()*?
- 3) What is the address of the third letter in the string *name[11]*?
- 4) What are the addresses (in the form register  $\pm$  offset) where the two static strings «*Insert your name:*» and «*You won!\n*» have been saved?
- 5) How many bytes are there between the starting of *name[11]* and the old base pointer saved on the stack?
- 6) What is the address (in the form register  $\pm$  offset) of the variable *n* given as argument to the function *fun()*?

- Can you execute the function *secret\_fun()* ?

# Thanks to...

